Machine Learning Assignment 1

Michelle Ho

CUSP-GX-5006
Assignment # 1

Abstract

In this assignment, I explore how the market value per square foot of buildings in Manhattan can be predicted by other variables for each building. The exercise demonstrates that the use of Principal Component Analysis (PCA) to alter the feature space prior to linear regression can help reduce noise and multicollinearity, leading to a smoother or more stable model.

1 Introduction

The main goal of this assignment is to demonstrate the use of linear regression on raw data, PCA-featured data, and finally a quick comparison to regularization methods, Lasso and Ridge regression.

This assignment makes use of a dataset on Manhattan buildings, provided by Professor Luis Gustavo Nonato. The dataset is somewhat problematic, with outliers and categorical variables mixed with non-categorical. In order to demonstrate the differences of the regression models, I chose to keep all the observations, even outliers. I also chose to treat some categorical variables like non-categorical at one point, even though this is incorrect and leads to very bad models. At a later step, I remove the categorical variables. The other goal of the assignment is to build a model that can predict market value per a square foot for buildings in Manhattan.

2 Methods and Data Sets

The provided dataset, 'manhattan-dof.csv', includes 2645 observations on 6 variables related to buildings in Manhattan, as described below:

- Neighborhood: a categorical variable for neighborhood. There are 38 unique neighborhoods in the dataset.
- BldClassif: a categorical variable describing the building classification. There are three unique classifications.
- YearBuilt: the year the building was built, ranges from 1864 to 2006.
- GrossSqFt: An integer for the gross square footage of the building.
- GrossIncomeSqFt: A float describing the rental value of a square foot of the building
- MarketValueperSqFt: A float describing the market value of a square foot of the building.
- The data was normalized by dividing by the maximum value for each variable.
- The dependent variable being predicted is 'MarketValueperSqFt' for all models (referred to as Y in this paper)

- The independent variables are all the others. At some steps, a subset of 'YearBuilt', 'GrossSqFt', and 'GrossIncomeSqFt' were used as the only independents.

The steps taken for this assignment:

1. First, an ordinary least squares linear regression model was fitted for the normalized, raw dataset (LM1).
2. Then, the dataset was projected into its principal components and the linear regression was re-run with all principal components (LM2).
3. The linear regression was re-run with the first three principal components (LM3).
4. The LM1 and LM3 linear models are assessed with cross validation.
5. An ordinary least squares linear regression model was fitted for the raw dataset with the categorical variables (BldClassif and Neighborhood) removed (LM4).
6. This new linear regression model is assessed with cross validation.
7. Finally, lasso and ridge regression are performed on the raw dataset with categorical variables removed.

3 Results

The first linear model generated was fitted for normalized, raw data with all dependent variables. The second linear model was fitted after projecting the data onto five principal components. The coefficients for these models are reported in Table 1. The x1, x2, x3, x4 and x5 coefficients have no interpretable meaning. The 'Neighborhood' and 'BldClassif' variables also have little interpretable value, since they are being treated as non-categorical at this point. Both models have the same Mean Square Error (MSE), as expected because I kept all five principal components, which can be thought of as a rotation of the "coordinate system" of our raw data. The MSE for both is about 0.02705.

The R-squared value for both models is very low (0.001), and every variable or component's p-value is above 0.08, indicating that these models are very poor. It's possible that these variables or components are zero or even the opposite sign. The plot of actual Y versus predicted Y for LM1 (Figure 1) shows that the model is predicting very close to 0.41 for almost all inputs, which is in fact the mean market value per square foot (without normalization, this is $128.75/square foot). In other words, this model is not very sophisticated and is only predicting the mean.

Next, I keep only the first three principal components. I chose the first three via the "elbow" method, where there appears to be a slight inflection point in the explained variance ratio (Figure 2). By incrementally increasing the number of principal components, I find that the mean square error approaches the mean square error of the regular linear model and that the percentage of explain variances approaches 100%. With three principal components, approximately 81% of the variance in the data is preserved.

In using PCA to project our raw data and then only selecting the first three principle components to fit the linear model, I am hoping to reduce the "noise" inside the data. This third linear model shows little improvement on the first two linear models. The R-squared is even lower, 0.0002, and three components have high p-values. The plot of actual vs predicted (Figure 3) shows even less sophistication. Remember, though, that I have not removed outliers and that PCA can be influenced by outliers pulling the components toward unhelpful directions. Furthermore, PCA-featured linear regression with fewer components than original variables always performs worse because there is a loss of information. The main goal is to reduce covariances, reduce dimensionality, and create a smoother model. In this situation, with so few independent variables to begin with, the loss of information may be doing more harm than good.

Nonetheless, I assess the performance of this PCA-featured linear model (LM3) and the first linear model (LM1) via cross validation for comparison. The data is split 33% to testing and the rest to training. To assess the first method, a linear model is fit on the training data and in-sample R-squared is

calculated. The fitted model is then used to predict the test data set, and the out-of-sample R-squared is calculated. Repeated 1000 times, this yields an in-sample R-squared of 0.0021 and out-of-sample R-squared of -0.0054.

Cross validation is also applied to the PCA-featured linear model (LM3). PCA is used only on the training set to find the principal components. The linear regression model is fit for these components. Then, the test data is projected onto the first three components, and the model predicts on the transformed test data. Repeated 1000 times, this yields an average in-sample R-squared of 0.00094 and out-of-sample R-squared of -0.00376.

It appears that the out-of-sample R-squared terms for both the PCA-featured and non-PCA featured linear models are both negative. The only way to achieve a negative R-squared is to have the predicted values error be worse than the error from simply predicting the mean. However, it is interesting to note that the in-sample R-squared for the LM1 model is slightly higher than the in-sample R-squared of the LM3 model, yet the out-of-sample R-squared of LM1 is slightly lower than LM3's. This may indicate that our PCA-featured linear model does a bit better out of sample, but not by much.

I repeat the steps for linear regression after removing the categorical variables 'Neighborhood' and 'BldClassif'. The results are reported in Table 4. Cross validation is also applied to this new model, and the average in-sample R-squared after 1000 repetitions is 0.0017 Out of Sample R-squared for 1000 times is -0.0035. Again, the model appears to be very poor, but there is slight improvement in the out-of-sample R-squared.

Finally, ridge and lasso regression models are also fitted using the non-categorical data. Figure 4 shows the predicted versus actual outcome of the L4 model (red), ridge regression model (green) and lasso regression model (black). Like the reduced dimension PCA-featured linear regression, these two shrinkage models also show the predicted values become simpler and smoother. Figure 5 shows the change in MSE as alpha, the regularization term, increases. Figure 6 shows how the coefficients shrink to zero as alpha increases.

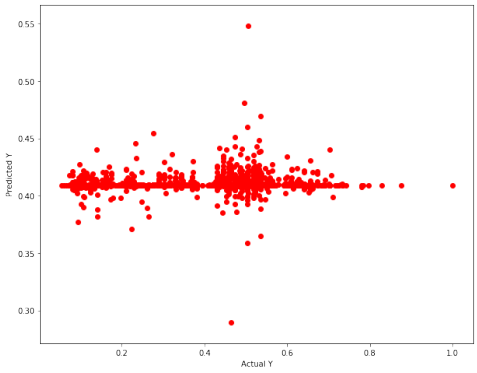

Figure 1: Actual vs Predicted – Simple Linear Model

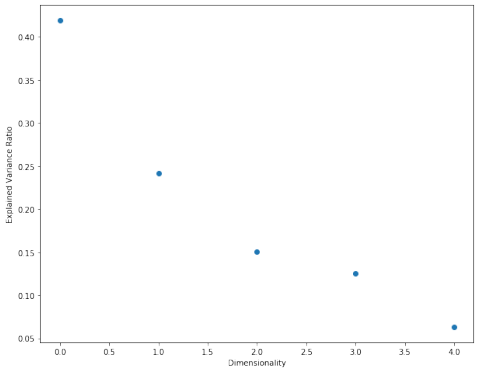

Figure 2: Explained Variance Ratio of Principal Components

Table 1: Results of Linear Models with Raw Normalized Data

	Normalized raw data (LM1)		PCA-featured data (LM2)
constant	0.40926343	constant	0.4101
Neighborhood	-0.12254491	x1	-0.0379
BldClassif	0.02690248	x2	-0.0113
YearBuilt	0.02695934	x3	-0.0441
GrossSqFt	0.13819696	x4	-0.1773
GrossIncomeSqFt	0.01101943	x5	-0.0273
R-squared	0.001	R-squared	0.001

Table 2: Results of Linear Model on Non-Categorical Data

	LM4
constant	0.4101
YearBuilt	-0.0626
GrossSqFt	0.1038
GrossIncomeSqFt	-0.1151
R-Squared	0.001

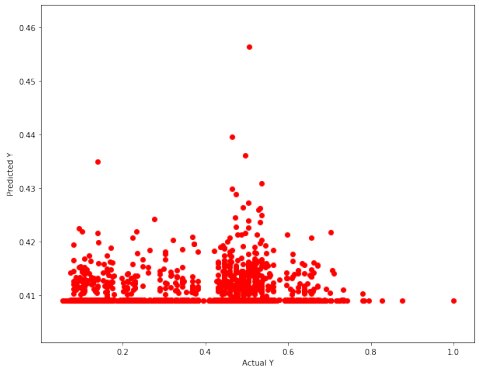

Figure 3: Actual vs Predicted – PCA-featured Linear Model

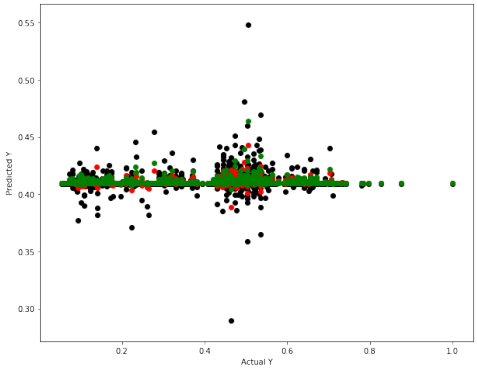

Figure 4: Actual vs Predicted – L4 model (red), ridge regression model (green) and lasso regression model (black)

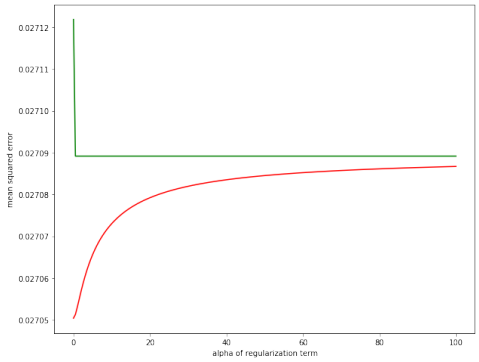

Figure 5: MSE for Ridge and Lasso, increasing alpha

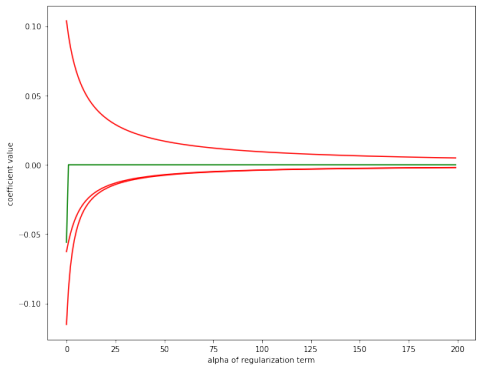

Figure 6: Coefficients for Ridge and Lasso, increasing alpha