Detecting Fraudulent Universities

Michelle Ho and Shay Lehmann

CUSP-GX-5006
Machine Learning Final Project

Abstract

For this project, the authors sought to examine higher education in the United States using the College Scorecard dataset. Their goals were twofold: 1) predict cohort's 3-year default rate based on information about the institution and 2) proactively identify potentially fraudulent schools.

1   Introduction

In the United States, the rising cost of a college degree and the corresponding increase in student loans is well-documented. Nonetheless, a college degree is the bare minimum requirement for many positions regardless of whether the specific major has provided any subject matter expertise relevant to the industry. The authors initially set out to identify the 'best value' colleges and chose the rate of student loan default on federally-funded loans as the metric with which to define value. Since most defaults are the result of a borrower being unable to make monthly payments, a default suggests that the student did not obtain a well-paying position. Furthermore, the default will negatively affect credit history for years to come. Exploratory analysis revealed significant discrepancies in the distribution of loan default by institution type (public, private not-for-profit and private for-profit), with public and private not-for-profits exhibiting similar patterns. Private for-profit schools, on the other hand, exhibited a significantly higher mean default rate as well as a much larger range of potential outcomes with some schools reaching default rates over 30% (Figure 1). The authors then began a secondary inquiry to determine whether they could identify fraudulent or potentially fraudulent universities. A number of for-profit universities, including the Corinthian College network have recently come under investigation by the federal government for their graduates' high default rates and misleading promises about job placement rates. In many instances, the federal government has forgiven several millions of dollars in outstanding loans. While loan forgiveness is one step towards aiding the victims, it does not compensate them or remediate inadequate education.

The authors hope machine learning techniques can be used to identify a school's 3-year loan default rate and fraudulent schools in particular. They hope that this will enable the government to take proactive steps towards investigating such universities before they are able to collect years of federal funds while delivering low quality education. They also hope that publicizing the high default rates will prompt reform of the process and build support for increased funding for public universities.

2   Methods and Data Sets

This report makes use of the College Scorecard dataset. This dataset is relatively new initiative to compile education-related from disparate federal institutions into one centralized location. Although the earliest cohorts

were from the late 1990s, measurements have changed over the years and many do not have data available for all variables. As such, the authors examined data from the 2011-2013 cohorts in order to predict the 3-year default rate for the cohort. This dataset includes 7414 schools and 1743 variables. However, there is not complete coverage for all variables depending on the institution. For this analysis, the data was subsetted based upon available features.

The variables are described below:

CONTROL: a categorical variable indicating the funding type of a school. Categories are 1="public", 2="private non-profit", or 3="private for-profit".

CDR3: Cohort Default Rate at 3 Years

MAIN: Boolean identifier for main campus

LO_INC_DEBT_MDN: The median debt for students with family income between $0-$30,000

HI_INC_DEBT_MDN: The median debt for students with family income $75,001+

MD_INC_DEBT_MDN: The median debt for students with family income between $30,001-$75,000

FIRST_GEN: Share of students who are first-generation students

COMP_ORIGIN_YR4_RT: Share of students who complete at the original university within 4 years

ADM_RATE_ALL: The admissions rate for the entire school

AGE_ENTRY: The average age at time of entry for students

INEXPFTE: Dollars spent by the institution on education per full-time student

The dependent variable being classified and predicted is 'CDR3' for the first step.

The data was limited to only the main campus for each institution as many features were reported as identical for each satellite campus leading to redundant entries. The MAIN column was then dropped for further analysis.

The independent variables are all the others.

Any observations with null values for any of the chosen variables are dropped. After these drops, the number of observations in the dataset is 2122 for the 2011 cohort, 2151 for the 2012 cohort, 2052 for the 2013 cohort.

Finally, the categorical variable CONTROL is binarized so that all samples are represented by boolean feature vectors.

The steps taken for this project were:

1. Random Forest Regression is run on the 2011 cohort using test sizes (as fraction of data) of .25, .33, .5, .67 and .75 with a depth of 8 and the resulting in- and out-of-sample R^2 visualized (Figure 2).

2. Random Forest Regression is run on the 2011 cohort using depths ranging from 3-11 and the resulting in- and out-of-sample R^2 visualized (Figure 3).

3. Random Forest Regressions are run on each cohort with a test size of .33, a max depth of 8 and 100 trees in the forest. Resulting actual v. predicted are visualized for both test and training samples (Figure 4, Figure 5, Figure 6)

4. Random Forest Regression is run using the 2011 cohort as training set and the 2013 cohort as test set (Figure 7).

5. Data from the 2011 and 2012 cohorts are grouped by school and averaged if the school was in both cohorts. Random Forest Regression is run on the averaged data and used to predict the 3-year default rate for the 2013 cohort (Figure 8)

3 Results

The two preliminary examinations of tree depth and test size were used to set the test size and max depth for the remainder of tests. In-sample R^2 was robust against test size, but out-of-sample R^2 was much more sensitive. With regard to the depths, in-sample R^2 and out-of-sample R^2 increased for both with more steps. However, out-of-sample plateaued at 8. Therefore, a test size of .33 and a maximum depth of 8 were used for subsequent models. Models were then trained within each cohort. The model trained on the 2011 cohort data had an in sample $R^2 = .928$ and an out-of-sample $R^2 = .802$. The most important features based on percentage of models in which they occurred

were FIRST_GEN (53.2%), MD_INC_DEBT_MDN (12.3%) and COMP_ORIG_YR4_RT (10.8%). The model trained on the 2012 cohort data The model trained on the 2011 cohort data had an in sample R^2 = .929 and an out-of-sample R^2 = .826. The most important features based on percentage of models in which they occurred were FIRST_GEN (55.7%) and COMP_ORIG_YR4_RT (20.5%). The model trained on the 2013 cohort data had an in sample R^2 = .915 and an out-of-sample R^2 = .772. The most important features based on percentage of models in which they occurred were COMP_ORIG_YR4_RT (46.1%), FIRST_GEN (17.9%) and CONTROL = Private For-Profit (15.8%). Since Private For-Profit institutions have such a wide range of default percentages as compared to the other two types, the tree often split on other factors that are concentrated within high default rates, namely a large share of first generation students and a low proportion who finish within four years.

Given the variety from year to year, the authors attempted to validate a model trained on 2011 data with results from 2013. However, this model lost some of its predictive power with a wider spread, especially in the highest default rates.

Lastly, the authors attempted to improve upon this year to year model by by averaging the features over years 2011 and 2012 in order to test 2013 data. This model did not improve upon the previous model and exhibited the same loss of predictive capability, particularly for the highest default rates.

The authors suspect that there may be additional factors that cause an exceptionally high default rate that may be missing from either the model or the dataset. Additional research into those schools which were greatly under-predicted may help to identify those factors.


Figure 1: Exploratory look at difference between default rates by Control Group, 2011

Figure 2: Exploration of Effect of Test Size on Random Forest Regression, 2011 Cohort

3

Figure 3: Exploration of Effect of Max Depth on Random Forest Regression, 2011 Cohort


Figure 5: Results of the model trained on data from the 2012 cohort

Figure 4: Results of the model trained on data from the 2011 cohort

Figure 6: Results of the model trained on data from the 2013 cohort

4

Figure 7: Results of the model trained on data from the 2011 cohort and tested on the 2013 cohort

Figure 8: Results of the model trained on averaged data from 2011-2012 cohorts and tested on the 2013 cohort